

Inverse Folding Models Sample Sequence Space Shaped by MHC Anchor Identity and TCR Context

Sergio E. Mares^{1,2}

¹Adimab, Mountain View, CA 94043, USA

²Center for Computational Biology, University of California, Berkeley, CA 94720, USA

Abstract

Inverse-folding models have emerged as potent tools for designing protein sequence from structure, and the degree to which they capture the biophysical constraints of a specific molecular interface stands as an open question. TCR cross-reactivity has caused fatal toxicity in the clinic, yet the peptide sequence space is far too large to screen experimentally. We therefore ask whether inverse-folding models can recover the native peptide from structure alone, and what their sampling distribution reveals about the constraints that give rise to cross-reactivity. We applied four inverse-folding models to fifty solved pMHC-TCR structures under two structural contexts, generating 3.9 million peptide designs. Every allele-matched design was scored against two independent binding predictors, and we characterized the full sampling distribution each model induces over peptide sequence space. No model recovered the native peptide. Recovery is instead bimodal across the two anchor positions that dominate MHC binding. The C-terminal anchor P Ω , though structurally load-bearing, is largely missed, while the position-2 anchor recovers at or above the interior average. Removing the T-cell receptor increases sampling diversity and raises predicted binding affinity for all four models, though the size of that affinity gain differs almost sixfold between them. It also lowers recovery at experimentally validated contact positions, and leaves every model that reports a confidence score less confident in its own designs. Both effects indicate that the models read the receptor as part of the interface. This suggests that inverse-folding models encode peptide compatibility as a set of position-specific structural constraints rather than a single optimal sequence. Our investigation establishes the sampling distribution as a probe of the constraints governing peptide recognition, and highlights both the promise and the limitations of structure-based models for predicting cross-reactivity.

1 Introduction

T-cell receptor (TCR) engineering therapies promise a transformative future for cancer treatment, and the first reached approval in 2024 [7, 12], yet cross-reactivity, the recognition of unintended peptide-MHC (pMHC) complexes, remains a fundamental barrier to scaling them into personalized medicine [20, 31]. An affinity-enhanced MAGE-A3

receptor that had cleared extensive preclinical testing ultimately caused fatal cardiac toxicity in two patients, cross-recognizing an unrelated pMHC derived from titin [5, 17].

Experimental methods therefore sample only a fraction of the pMHC-TCR interaction space, which is combinatorially vast, roughly 20^L for a peptide of length L , yet far from uniform. The anchor sites collapse the entropy at fixed positions while leaving the remainder variable [11, 23], so the effective dimensionality of the presented distribution falls far below its combinatorial ceiling. It is this low-dimensional structure that makes the problem learnable, yet the same data limitations that bound experimental sampling also limit the models trained on that data [22, 24], motivating approaches that exploit structural rather than purely sequence-based signals.

Inverse-folding models sample sequences from a distribution conditioned on structure [8, 13], and on the receptor side of this interface, the approach has recent precedent. Comparing ProteinMPNN and ESM-IF1 against physics-based Rosetta design [1] across 38 solved TCR-pMHC complexes, prior work [25] found that both inverse-folding models recover native CDR3 interface residues well above chance, establishing that inverse folding captures real, non-trivial signals from a TCR-pMHC interface.

Here, we turn the same class of models toward the peptide itself. The point estimate asks whether the model recovers the native peptide the structure was solved with; the distribution asks which other peptides the model treats as compatible with the same groove and the same receptor, precisely the variant and off-target sequences that underlie cross-reactivity. Because inverse-folding models score structural compatibility rather than a labeled binding phenotype, that distribution provides a first-pass map of the tolerated peptide neighborhood around the native ligand, and, because MHC-groove geometry and TCR-contact geometry occupy structurally distinguishable positions along the peptide, acts as a probe of which positions the model treats as side-chain-constrained versus permissive. We therefore ask whether inverse-folding models recover the native peptide, and whether the distribution they induce over sequence space concentrates on the tolerated neighborhood underlying cross-reactivity and off-target recognition.

2 Materials and Methods

2.1 Panel and models

Fifty pMHC-TCR crystal structures were retrieved from the Protein Data Bank [2] (Supplementary Table S1). Every structure includes a bound peptide, the MHC chains, and a complete TCR $\alpha\beta$ heterodimer [3] resolved in the same asymmetric unit with peptides of various lengths (Supplementary Figure S1a); the class I entries additionally include β_2 -microglobulin where it was resolved. Forty structures are HLA-A*02:01. The remaining ten are one HLA-B*51:01 (4MJJ), one HLA-B*08:01, one HLA-B*35, two mouse H-2L^d, and five MHC class II complexes, all of which were excluded from the allele-specific binding predictors. The average crystallographic resolution is 2.51 (Supplementary Figure S1b).

The panel is not fifty independent observations. It contains 33 distinct peptides and 26 distinct CDR3 β sequences, with SLLMWITQC solved in five crystals and several receptors (A6, 1G4, DMF5) recurring, so a pooled statistic weights an epitope by how often it happens to have been crystallised. We therefore also report a non-redundant set: human MHC class I only, one structure per TCR family (CDR3 β clustered so clonotypes differing by one or two substitutions collapse together), and one structure per peptide, taking the best-resolution member of each group as its representative so the choice is independent of any score. This leaves 30 structures over 26 peptides when further restricted to HLA-A*02:01 (Supplementary Figure S17).

The four models are ProteinMPNN [8], a message-passing GNN encoding the fixed backbone as a k -nearest-neighbor graph and decoding sequence autoregressively; ProteinMPNN (no MHC), the same architecture retrained with MHC/TCR-class structures excluded; ESM-IF1 [13] and LigandMPNN [9] (Supplementary Table S2).

Each of the four tools samples only the peptide chain, with the MHC, β_2 -microglobulin, and (where present) TCR chains held fixed as structural context. Two context conditions were run for every structure, MHC + β_2 m + TCR $\alpha\beta$ and the TCR removed from the input structure entirely. Each condition was sampled to 10,000 raw designs at a fixed sampling temperature of $T = 0.1$.

3 Results

3.1 Models Sample Strong Position-Specific Sequence Bias

To establish how much of the peptide sequence space these models actually explore, we sampled 9,984 designs per (structure, model) cell in the pMHC+TCR context and counted how many were distinct. An average of only 79 were unique, 0.81% of the draws, ranging from 48 for ESM-IF1 to 105 for ProteinMPNN (no MHC) with 89 for ProteinMPNN and 72 for LigandMPNN, so these models return redundancy rather than diversity (Supplementary Figure S2). Sequence logos [27] show every model imposes strong position-specific restriction independent of recovery (Supplementary Figures S3–S4d), a property of the sampling distribution rather than evolutionary conservation, tracing the

same contacts seen in the crystals, where 29 of 82 MHC positions contacting the peptide anywhere in the panel do so in at least 40 of the 50 (Supplementary Figure S1c).

3.2 The two MHC anchors impose different constraints and recover differently

To determine which peptide positions carry recoverable signal, we measured per-position recovery of each crystal's own native residue, pooled across fifty structures and four models. Recovery is bimodal rather than uniform, and the two positions where this matters most are the two positions immunology already identifies as anchors [11, 18] (Figure 1). P2, seated in pocket B, constrains a specific side chain and recovers well above the interior average (0.65 vs. 0.50), whereas P Ω , seated in pocket F, imposes a permissive backbone constraint and recovers significantly *below* it (0.26 vs. 0.50, $p = 0.00021$). Both are structurally indispensable, yet only P2 gives a design model a signal it can exploit. Neither result is an artefact of the panel's redundancy. On the non-redundant HLA-A*02:01 set (30 structures, 26 peptides, one per TCR family) P2 still recovers above interior (0.64 vs. 0.52, $p = 0.010$) and P Ω still falls below it (0.17, $p < 10^{-4}$), and the P Ω deficit in fact deepens as repeated epitopes are removed, from 0.25 on the full panel to 0.17 here (Supplementary Figure S17).

To ask whether that asymmetry reflects side-chain identity alone or a coarser physicochemical preference, we compared each native anchor residue against its own peptide's interior positions on the Kyte-Doolittle hydrophobicity scale [16], and separately scored every design for whether it preserved the native residue's side-chain class (Supplementary Figure S12). Both anchors are markedly more hydrophobic than the interior they sit alongside (3.38 at P2 and 3.80 at P Ω , against an interior mean of 0.86; $p < 0.0001$ for both, paired Wilcoxon across 40 structures), confirming that both pockets recruit hydrophobic residues. The two anchors nonetheless part company on how tightly that preference is enforced in the designs. P2 designs preserve the native side-chain class well above the interior baseline (86.1% vs. 70.8%, $p = 0.0002$), whereas P Ω designs do not (67.2% vs. 70.8%, $p = n.s.$), even though P Ω designs do show significantly reduced hydrophobicity variance relative to the interior ($p = 0.0032$). P Ω therefore behaves as a pocket that biases toward a hydrophobic character without fixing which hydrophobic residue satisfies it, which is what a low exact-identity recovery at a structurally load-bearing position should look like.

Averaged across the entire peptide, the four models fall within a narrow band of mean recovery (0.43 to 0.47 in full context and 0.29 to 0.36 without the TCR; Supplementary Figure S5), narrow enough that the models might appear indistinguishable by inspection. A Friedman test indicates otherwise in both contexts ($\chi^2 = 13.13$, $p = 0.004$, full context; $\chi^2 = 19.38$, $p < 0.001$, mhonly), so this apparent similarity reflects proximity in magnitude rather than statistical equivalence. As an aside, on 3QDG the central Gly-Ile-Gly motif at positions 4–6, well characterized in the literature [4], recovers essentially perfectly here (1.00, 0.73, 1.00 vs. 0.50), though we have not attempted a systematic motif

9-mers -- pMHC+TCR, n=32 structures x 4 models (violin = distribution shape, dots = every raw point; labels = % of points at $\geq 99\%$ and at $\leq 1\%$ recovery)




Figure 1: Per-position recovery is bimodal, and the two MHC anchors diverge. Recovery of the native residue across the 32 nonamer crystals and 4 models, one panel per peptide position (pMHC+TCR context). The violin gives the distribution shape and every (crystal, model) value is overlaid as a point. Recovery is sharply bimodal at most positions, piling up near 0 and 1 rather than spreading around a mean. Positions are only comparable within a peptide length, so each length is shown separately; the remaining lengths are in Supplementary Figures S13 and S14.

annotation across the full panel.

The relationship between resolution and recovery illustrates how pooling across a heterogeneous panel can understate a genuine effect. Pooled across all fifty structures, the correlation is modest, and only one model survives correction individually ($r = -0.29$ pooled, $p = 0.040$; ESM-IF1 alone, $r = -0.37$, $p = 0.008$; Supplementary Figure S6). However, five of the fifty structures (2P5W, 2BNR, 2F53, 2F54 and 2PYE) are independent crystals of the identical peptide (SLLMWITQC, the NY-ESO-1 epitope [6, 26]) differing primarily in resolution (1.9–2.7 Å). Holding peptide identity fixed in this way removes a confound the pooled panel cannot. Declining resolution drives recovery down for ProteinMPNN (no MHC) ($r = -0.97$, $p = 0.005$) and, less clearly, ESM-IF1 ($r = -0.84$, $p = 0.075$), and drives unique design counts up for both ProteinMPNN variants ($r = 0.97$, $p = 0.007$ and $r = 0.95$, $p = 0.012$; Supplementary Figure S7). It is not universal: LigandMPNN moves in the opposite direction on both measures ($r = +0.19$ and $r = -0.20$, both n.s.), so the pooled-panel estimate is diluted by heterogeneity rather than simply attenuated.

3.3 Most peptides sampled by inverse-folding models are predicted to be presented

Having established the recovery rates of the native peptide and the sequence space each structure samples, we next asked whether those sampled peptides would actually be presented, scoring every unique design from the forty HLA-A*02:01 structures for predicted HLA-A2 binding with two unrelated predictors, MHCflurry [22] and ESM-CBA [10, 19], neither of which had seen this panel (Figure 2).

The two predictors agree closely on the same designs




Figure 2: MHCflurry and ESM-CBA predictions of designs $\log_{10}(\text{predicted IC}_{50}, \text{nM})$ with pMHC+TCR context, one panel per model. Dashed lines mark the 500 nM; The percentage is the fraction of designs under 500 nM by both predictors at once.

($r = 0.83$ – 0.93 per model in full context, up to 0.94 without the TCR), and agreement between two methodologically unrelated predictors constitutes stronger evidence than either predictor alone. On this shared landscape, projected jointly into one UMAP embedding space [21] (Supplementary Figure S11), ProteinMPNN and LigandMPNN consis-

tently sample toward higher predicted affinity than ESM-IF1. At the 500 nM cutoff [28], 52–64% of ProteinMPNN-family and LigandMPNN designs qualify by MHCflurry, compared with 41% for ESM-IF1 ($\chi^2 = 238.50$, $p = 2 \times 10^{-51}$; see also Section 3.4). This pattern is better characterized as ESM-IF1 being a consistent low-affinity outlier than as a general superiority of ProteinMPNN over ESM-IF1, since the other three tools do not share this deficit.

3.4 Without the TCR, inverse-folding models sample more widely but less confidently

To understand the effect of the TCR on what these models propose, we resampled every structure with the TCR deleted from the input and compared the resulting designs against the full-complex designs above.

Every model explores substantially more sequence space without the TCR, with unique design counts roughly doubling or tripling (ProteinMPNN 2047 $\rightarrow$ 6637; ProteinMPNN (no MHC) 2363 $\rightarrow$ 4478; ESM-IF1 1783 $\rightarrow$ 3459; LigandMPNN 1635 $\rightarrow$ 4426; Supplementary Figure S5). This larger space skews toward higher predicted affinity, but not uniformly across models. ProteinMPNN (no MHC) shows the largest shift (strong-binder rate 60.3% $\rightarrow$ 86.0%), LigandMPNN a smaller but genuine one (52.2% $\rightarrow$ 62.6%), while ProteinMPNN and ESM-IF1 show little change (Supplementary Figure S10). The affinity gain associated with TCR removal thus appears concentrated in the two models already most context-sensitive, rather than reflecting a general consequence of removing the TCR.

A related but distinct signal comes from each model's own confidence in the sequences it designs, independent of either external predictor. ProteinMPNN and ProteinMPNN (no MHC) report an average per-residue negative log-likelihood for each design, for which lower values indicate higher confidence; LigandMPNN reports an overall confidence score, for which higher values indicate higher confidence. Pooled across all forty HLA-A*02:01 structures, removing the TCR shifts this internal score for every model, and in every case toward lower self-assessed confidence (ProteinMPNN mean 1.016 $\rightarrow$ 1.295; ProteinMPNN (no MHC) 0.925 $\rightarrow$ 1.128; LigandMPNN 0.413 $\rightarrow$ 0.321; Mann-Whitney $p < 10^{-200}$ for all three; Supplementary Figure S9a). This shift is not uniform across the panel. Broken out per structure and Bonferroni-corrected, it reaches significance in 15–17 of 19 structures depending on the model, but not in the remaining two to four (Supplementary Figure S9b). That a model can become less confident in its own designs by its own internal metric while simultaneously being scored as a stronger binder by an external predictor that never saw those designs indicates that the two signals are not redundant, since one reflects how comfortably a sequence fits the model's learned distribution, the other how likely it is to bind, and removing the TCR does not move them the same way. This distinction is more informative than amino-acid-level recovery or composition alone, since it speaks to what the model itself, not just an external predictor, "believes" about the sequences it is generating.

3.5 Recovery Loss Concentrates at Experimentally Validated TCR-Contact Positions

To understand how these residue interactions contribute functionally to TCR binding, we tested whether the recovery lost without the TCR concentrates at positions that experimental data already implicate in receptor engagement. We cross-referenced SKEMPI 2.0 [14] and IEDB [30] against each structure's crystallized TCR clone (matched by CDR3 sequence) and identified 21 (structure, position) pairs with an experimentally measured mutation that alters TCR binding. Pooling recovery across exactly these positions, TCR-present designs recover the native residue 54–58% of the time, compared with 34–38% for TCR-absent designs (Figure 3), a consistent, though not universal, gap; a minority of individual positions move in the opposite direction.

That comparison rests on 21 curated positions. The extended panel supports a more direct version of the same test, against measured binding energies rather than curated annotations. Restricting SKEMPI to single mutations inside the epitope itself yields 55 (complex, position) pairs across 17 complexes with a measured $\Delta\Delta G$. At positions where mutation costs at least 1 kcal/mol, the recovery gain from having the TCR present is +0.150, against +0.018 elsewhere ($p = 0.047$, Mann-Whitney; Supplementary Figure S15). With the TCR deleted the distinction disappears outright (0.243 vs. 0.172, $p = 0.85$), which is the null this account predicts. We report the direction rather than the effect size: critical positions lead in every subset we examined, but the pooled p rises to 0.19–0.22 when the three same-epitope H-2A crystals or the class II entries are excluded, and survives only 6 of 17 leave-one-complex-out folds. Half the positions come from a single epitope, so this is a consistent direction on an underpowered sample rather than an established effect.

A second, independent line of evidence separates what these scores do encode from what they do not. 6AM5 and 6AMU present the same DMF5 receptor and the same HLA-A*02:01 groove with two different peptides, so threading one peptide set onto both backbones isolates conformation from sequence. Across all three scoring models, each structure's own crystallized peptide ranks 2.3 of 13 on its own backbone and 11.0 of 13 on the other (Figure 4). The score therefore reads the backbone it is conditioned on, not the peptide sequence alone. Measured affinity behaves differently. The same 13 peptides carry published K_D and thermal-stability values, nine of them as an actual measurement rather than a bound. Across those nine the score carries no information about affinity on either backbone, for any of the three models ($|\rho| \leq 0.32$, all $p > 0.4$; Supplementary Figure S16). Including the four peptides reported only as $> 500 \mu\text{M}$ or as not detected raises the apparent correlation to about ± 0.5 and reverses its sign between the two backbones, but that follows from where the assay ceiling falls relative to the two peptide families rather than from the scores, which is why those peptides are excluded. The distinction matters for how these models should be used: they report which peptides are compatible with a given groove




Figure 3: **Where the TCR-context benefit actually concentrates.** Recovery(full) – recovery(mhconly), per crystal per model, crystals ordered by resolution (best to worst, left to right).




Figure 4: **Each backbone favours its own crystallized peptide.** The same 13 DMF5 peptides threaded onto 6AM5 (blue) and 6AMU (orange) and ordered by their score on 6AM5, shown as the mean of three models with a ± 1 SD band, with rings and dotted lines marking each structure’s own peptide.

and backbone, the quantity relevant to cross-reactivity, and not a ranking of binding strength.

4 Discussion

Inverse-folding models have rapidly advanced structure-based protein design [8, 9, 13], and their ability to recover

native interface residues above chance [25] invites the question of whether they capture the biophysics of an interface or merely reproduce its most constrained positions. High recovery is easily read as the former. Our results point to the latter, since knowing that MHC anchor positions matter explains little on its own. P2 and P Ω are both structurally indispensable yet behave oppositely under inverse folding, one a side-chain readout, the other a backbone constraint tolerant of many side chains. This finer distinction is what a sequence-design model can meaningfully report. The same caution against pooling applies to resolution, where a weak correlation across fifty structures conceals a model-specific effect once peptide identity is controlled for, using the five SLLMWITQC/NY-ESO-1 crystals as a same-peptide replicate set. TCR removal resists a single verdict for the same reason. Diversity and predicted affinity both increase for every model, but by margins that differ almost sixfold between them, and recovery decreases on average without decreasing everywhere, because whether the model requires the TCR was never one question.

References

- [1] Rebecca F. Alford, Andrew Leaver-Fay, Jeliasko R. Jeliaskov, Matthew J. O’Meara, Frank P. DiMaio, Hahnbeom Park, Maxim V. Shapovalov, P. Douglas Renfrew, Vikram K. Mulligan, Kalli Kappel, Jason W. Labonte, Michael S. Pacella, Richard Bonneau, Philip Bradley, Roland L. Dunbrack, Rhiju Das, David Baker, Brian Kuhlman, Tanja Kortemme, and Jeffrey J. Gray. The Rosetta all-atom energy function for macromolecular modeling and design. *Journal of Chemical Theory and Computation*, 13(6):3031–3048, 2017. doi: 10.1021/acs.jctc.7b00125.

- [2] Helen M. Berman, John Westbrook, Zukang Feng, Gary Gilliland, T. N. Bhat, Helge Weissig, Ilya N. Shindyalov, and Philip E. Bourne. The Protein Data Bank. *Nucleic Acids Research*, 28(1):235–242, 2000. doi: 10.1093/nar/28.1.235.
- [3] Pamela J. Bjorkman, Mark A. Saper, Boudjema Samraoui, William S. Bennett, Jack L. Strominger, and Don C. Wiley. Structure of the human class I histocompatibility antigen, HLA-A2. *Nature*, 329(6139):506–512, 1987. doi: 10.1038/329506a0.
- [4] Oleg Y. Borbulevych, Sujatha M. Santhanagopalan, Moushumi Hossain, and Brian M. Baker. TCRs used in cancer gene therapy cross-react with MART-1/Melan-A tumor antigens via distinct mechanisms. *The Journal of Immunology*, 187(5):2453–2463, 2011. doi: 10.4049/jimmunol.1101268.
- [5] Brian J. Cameron, Andrew B. Gerry, Joseph Dukes, Jane V. Harper, Vivekanandan Kannan, Frayne C. Bianchi, Francis Grand, Joanna E. Brewer, Minnal Gupta, Gabriela Plesa, Giovanna Bossi, Annelise Vuidepot, Alex S. Powlesland, Alison Legg, Katherine J. Adams, Alan D. Bennett, Nicholas J. Pumphrey, Daniel D. Williams, Gwendolyn Binder-Scholl, Irina Kulikovskaya, Bruce L. Levine, James L. Riley, Angel Varela-Rohena, Edward A. Stadtmauer, Aaron P. Rapoport, Gerald P. Linette, Carl H. June, Namir J. Hassan, Michael Kalos, and Bent K. Jakobsen. Identification of a titin-derived HLA-A1-presented peptide as a cross-reactive target for engineered MAGE A3-directed T cells. *Science Translational Medicine*, 5(197):197ra103, 2013. doi: 10.1126/scitranslmed.3006034.
- [6] Jean-Li Chen, Guillaume Stewart-Jones, Giovanna Bossi, Nikolai M. Lissin, Linda Wooldridge, Eric Man Lik Choi, Gerhard Held, P. Rod Dunbar, Robert M. Esnouf, Malkit Sami, Jonathan M. Boulter, Pierre Rizkallah, Christoph Renner, Andrew Sewell, P. Anton van der Merwe, Bent K. Jakobsen, Gillian Griffiths, E. Yvonne Jones, and Vincenzo Cerundolo. Structural and kinetic basis for heightened immunogenicity of T cell vaccines. *Journal of Experimental Medicine*, 201(8):1243–1255, 2005. doi: 10.1084/jem.20042323.
- [7] Y. Nancy Chen and Anna M. Bulek. Approval of the first TCR-based cell therapy. *Molecular Therapy*, 32(10):3241–3242, 2024. doi: 10.1016/j.ymthe.2024.09.007. Afamitresgene autoleucel (TECELRA) received U.S. FDA accelerated approval on 2 August 2024.
- [8] Justas Dauparas, Ivan Anishchenko, Nathaniel Bennett, Hua Bai, Robert J. Ragotte, Lukas F. Milles, Basile I. M. Wicky, Alexis Courbet, Rob J. de Haas, Neville Bethel, Philip J. Y. Leung, Timothy F. Huddy, Sam Pellock, Doug Tischer, Frederick Chan, Bruce Koepnick, Hannah Nguyen, Alex Kang, Banumathi Sankaran, Asim K. Bera, Neil P. King, and David Baker. Robust deep learning-based protein sequence design using ProteinMPNN. *Science*, 378(6615):49–56, 2022. doi: 10.1126/science.add2187.
- [9] Justas Dauparas, Gyu Rie Lee, Robert Pecoraro, Lulu An, Ivan Anishchenko, Cameron Glasscock, and David Baker. Atomic context-conditioned protein sequence design using LigandMPNN. *Nature Methods*, 22(4):717–723, 2025. doi: 10.1038/s41592-025-02626-1.
- [10] EvolutionaryScale Team. ESM Cambrian: revealing the mysteries of proteins with unsupervised learning. Technical report, EvolutionaryScale, 2024. <https://evolutionaryscale.ai/blog/esm-cambrian>.
- [11] Kirsten Falk, Olaf Rötzschke, Stefan Stevanović, Günther Jung, and Hans-Georg Rammensee. Allele-specific motifs revealed by sequencing of self-peptides eluted from MHC molecules. *Nature*, 351(6324):290–296, 1991. doi: 10.1038/351290a0.
- [12] David S. Hong, Brian A. Van Tine, Sarina Biswas, Cheryl McAlpine, Melissa L. Johnson, Anthony J. Olaszanski, Jeffrey M. Clarke, Dejka Araujo, George R. Blumenschein, Partow Kebriaei, Quan Lin, Alex J. Tipping, Joseph P. Sanderson, Ruoxi Wang, Trupti Trivedi, Stavros Rafail, Nancy Rafail, Elliot Norry, and Aaron P. Rapoport. Autologous T cell therapy for MAGE-A4⁺ solid tumors in adults with advanced disease: a phase 1 trial. *Nature Medicine*, 29(1):104–114, 2023. doi: 10.1038/s41591-022-02128-z.
- [13] Chloe Hsu, Robert Verkuil, Jason Liu, Zeming Lin, Brian Hie, Tom Sercu, Adam Lerer, and Alexander Rives. Learning inverse folding from millions of predicted structures. In *Proceedings of the 39th International Conference on Machine Learning (ICML)*, volume 162 of *PMLR*, pages 8946–8970, 2022.
- [14] Justina Jankauskaitė, Brian Jiménez-García, Justas Dapkūnas, Juan Fernández-Recio, and Iain H. Moal. SKEMPI 2.0: an updated benchmark of changes in protein-protein binding energy, kinetics and thermodynamics upon mutation. *Bioinformatics*, 35(3):462–469, 2019. doi: 10.1093/bioinformatics/bty635.
- [15] John Jumper, Richard Evans, Alexander Pritzel, Tim Green, Michael Figurnov, Olaf Ronneberger, Kathryn Tunyasuvunakool, Russ Bates, Augustin Žídek, Anna Potapenko, Alex Bridgland, Clemens Meyer, Simon A. A. Kohl, Andrew J. Ballard, Andrew Cowie, Bernardino Romera-Paredes, Stanislav Nikolov, Rishub Jain, Jonas Adler, Trevor Back, Stig Petersen, David Reiman, Ellen Clancy, Michal Zielinski, Martin Steinegger, Michalina Pacholska, Tamas Berghammer, Sebastian Bodenstein, David Silver, Oriol Vinyals, Andrew W. Senior, Koray Kavukcuoglu, Pushmeet Kohli, and Demis Hassabis. Highly accurate protein structure prediction with AlphaFold. *Nature*, 596(7873):583–589, 2021. doi: 10.1038/s41586-021-03819-2.

- [16] Jack Kyte and Russell F. Doolittle. A simple method for displaying the hydropathic character of a protein. *Journal of Molecular Biology*, 157(1):105–132, 1982. doi: 10.1016/0022-2836(82)90515-0.
- [17] Gerald P. Linette, Edward A. Stadtmauer, Marcela V. Maus, Aaron P. Rapoport, Bruce L. Levine, Lyndsey Emery, Leslie Litzky, Adam Bagg, Beatriz M. Carreno, Patrick J. Cimino, Gwendolyn K. Binder-Scholl, Dominic P. Smethurst, Andrew B. Gerry, Nicholas J. Pumphrey, Alan D. Bennett, Joanna E. Brewer, Joseph Dukes, Jane Harper, Helen K. Tayton-Martin, Bent K. Jakobsen, Namir J. Hassan, Michael Kalos, and Carl H. June. Cardiovascular toxicity and titin cross-reactivity of affinity-enhanced T cells in myeloma and melanoma. *Blood*, 122(6):863–871, 2013. doi: 10.1182/blood-2013-03-490565.
- [18] Dean R. Madden. The three-dimensional structure of peptide-MHC complexes. *Annual Review of Immunology*, 13:587–622, 1995. doi: 10.1146/annurev.iy.13.040195.003103.
- [19] Sergio Mares, Ariel Espinoza, and Nilah Ioannidis. Continued domain-specific pre-training of protein language models for pMHC-I binding prediction. In *Proceedings of the 20th Machine Learning in Computational Biology Meeting*, volume 311 of *PMLR*, pages 304–325, 2025.
- [20] Don Mason. A very high level of crossreactivity is an essential feature of the T-cell receptor. *Immunology Today*, 19(9):395–404, 1998. doi: 10.1016/S0167-5699(98)01299-7.
- [21] Leland McInnes, John Healy, and James Melville. UMAP: uniform manifold approximation and projection for dimension reduction. *arXiv preprint arXiv:1802.03426*, 2018. doi: 10.48550/arXiv.1802.03426.
- [22] Timothy J. O’Donnell, Alex Rubinsteyn, and Uri Laserson. MHCflurry 2.0: improved pan-allele prediction of MHC class I-presented peptides by incorporating antigen processing. *Cell Systems*, 11(1):42–48.e7, 2020. doi: 10.1016/j.cels.2020.06.010.
- [23] Hans-Georg Rammensee, Jutta Bachmann, Niels P. N. Emmerich, Oskar A. Bachor, and Stefan Stevanović. SYFPEITHI: database for MHC ligands and peptide motifs. *Immunogenetics*, 50(3–4):213–219, 1999. doi: 10.1007/s002510050595.
- [24] Birkir Reynisson, Bruno Alvarez, Sinu Paul, Björn Peters, and Morten Nielsen. NetMHCpan-4.1 and NetMHCIIpan-4.0: improved predictions of MHC antigen presentation by concurrent motif deconvolution and integration of MS MHC eluted ligand data. *Nucleic Acids Research*, 48(W1):W449–W454, 2020. doi: 10.1093/nar/gkaa379.
- [25] Helder V. Ribeiro-Filho, Gabriel E. Jara, João V. S. Guerra, Melina Cheung, Nathaniel R. Felbinger, José G. C. Pereira, Brian G. Pierce, and Paulo S. Lopes-de Oliveira. Exploring the potential of structure-based deep learning approaches for T cell receptor design. *PLOS Computational Biology*, 20(10):e1012489, 2024. doi: 10.1371/journal.pcbi.1012489.
- [26] Malkit Sami, Pierre J. Rizkallah, Steven Dunn, Peter Molloy, Ruth Moysey, Annelise Vuidepot, Emma Baston, Penio Todorov, Yi Li, Fang Gao, Jonathan M. Boulter, and Bent K. Jakobsen. Crystal structures of high affinity human T-cell receptors bound to peptide major histocompatibility complex reveal native diagonal binding geometry. *Protein Engineering, Design and Selection*, 20(8):397–403, 2007. doi: 10.1093/protein/gzm033.
- [27] Thomas D. Schneider and R. Michael Stephens. Sequence logos: a new way to display consensus sequences. *Nucleic Acids Research*, 18(20):6097–6100, 1990. doi: 10.1093/nar/18.20.6097.
- [28] Alessandro Sette, Antonella Vitiello, Barbara Rehman, Peggy Fowler, Ramila Nayersina, W. Martin Kast, Cornelis J. M. Melief, Carla Oseroff, Lauri Yuan, John Ruppert, John Sidney, Maria-França del Guercio, Scott Southwood, Ralph T. Kubo, Robert W. Chesnut, Howard M. Grey, and Francis V. Chisari. The relationship between class I binding affinity and immunogenicity of potential cytotoxic T lymphocyte epitopes. *The Journal of Immunology*, 153(12):5586–5592, 1994. doi: 10.4049/jimmunol.153.12.5586.
- [29] Ian Sillitoe, Nicola Bordin, Natalie Dawson, Vaishali P. Waman, Paul Ashford, Harry M. Scholes, Camilla S. M. Pang, Laurel Woodridge, Clemens Rauer, Nee-ladri Sen, Mahnaz Abbasian, Sean Le Cornu, Su Datta Lam, Karel Berka, Ivana Hutarova Varekova, Radka Svobodova, Jon Lees, and Christine A. Orengo. CATH: increased structural coverage of functional space. *Nucleic Acids Research*, 49(D1):D266–D273, 2021. doi: 10.1093/nar/gkaa1079.
- [30] Randi Vita, Swapnil Mahajan, James A. Overton, Sandeep Kumar Dhanda, Sheridan Martini, Jason R. Cantrell, Daniel K. Wheeler, Alessandro Sette, and Bjoern Peters. The Immune Epitope Database (IEDB): 2018 update. *Nucleic Acids Research*, 47(D1):D339–D343, 2019. doi: 10.1093/nar/gky1006.
- [31] Linda Wooldridge, Julia Ekeruche-Makinde, Hugo A. van den Berg, Anna Skowera, John J. Miles, Mai Ping Tan, Garry Dolton, Mathew Clement, Sian Llewellyn-Lacey, David A. Price, Mark Peakman, and Andrew K. Sewell. A single autoimmune T cell receptor recognizes more than a million different peptides. *Journal of Biological Chemistry*, 287(2):1168–1177, 2012. doi: 10.1074/jbc.M111.289488.

A Supplementary Figures

PDB	Peptide	Length	Allele	Resolution (Å)
2P5W	SLLMWITQC	9	HLA-A*02:01	2.20
1QSF	LLFGYPVAV	9	HLA-A*02:01	2.80
1QRN	LLFGYAVYV	9	HLA-A*02:01	2.80
2BNR	SLLMWITQC	9	HLA-A*02:01	1.90
2GJ6	LLFGKPVYV	9	HLA-A*02:01	2.56
2F53	SLLMWITQC	9	HLA-A*02:01	2.10
2F54	SLLMWITQC	9	HLA-A*02:01	2.70
3QDG	ELAGIGILTV	10	HLA-A*02:01	2.69
3QEQ	AAGIGILTV	9	HLA-A*02:01	2.59
3QFJ	LLFGFPVYV	9	HLA-A*02:01	2.29
3GSN	NLVPMVATV	9	HLA-A*02:01	2.80
1OGA	GILGFVFTL	9	HLA-A*02:01	1.40
3UTS	ALWGPDPAAA	10	HLA-A*02:01	2.71
5C0A	MVWGPDPLYV	10	HLA-A*02:01	2.46
5C0B	RQFGPDFPTI	10	HLA-A*02:01	2.03
5HHO	GILEFVFTL	9	HLA-A*02:01	2.95
5EU6	YLEPGPVTV	9	HLA-A*02:01	2.02
2VLR	GILGFVFTL	9	HLA-A*02:01	2.30
4MJI	TAFTIPSI	8	HLA-B*51:01	2.99
5NME	SLYNTVATL	9	HLA-A*02:01	2.94
1BD2	LLFGYPVYV	9	HLA-A*02:01	2.50
1LP9	ALWGFFPVL	9	HLA-A*02:01	2.00
1MI5	FLRGRAYGL	9	HLA-B*08:01	2.50
1QSE	LLFGYPRYV	9	HLA-A*02:01	2.80
2AK4	LPEPLPQGQLTAY	13	HLA-B*35	2.50
2BNQ	SLLMWITQV	9	HLA-A*02:01	1.70
2E7L	QLSPFPFDL	9	H-2L ^d (mouse)	2.50
2J8U	ALWGFFPVL	9	HLA-A*02:01	2.88
2JCC	ALWGFFPVL	9	HLA-A*02:01	2.50
2OI9	QLSPFPFDL	9	H-2L ^d (mouse)	2.35
2PYE	SLLMWITQC	9	HLA-A*02:01	2.30
2UWE	ALWGFFPVL	9	HLA-A*02:01	2.40
3C60	FEAQKAKANKAVD	13	H-2A (mouse, class II)	3.05
3D3V	LLFGFPVYV	9	HLA-A*02:01	2.80
3H9S	MLWGYLQYV	9	HLA-A*02:01	2.70
3PWP	LGYGFNVI	9	HLA-A*02:01	2.69
3QDJ	AAGIGILTV	9	HLA-A*02:01	2.30
3QIB	ADLIAYLKQATKG	13	I-E ^k (mouse, class II)	2.70
4FTV	LLFGYPVYV	9	HLA-A*02:01	2.74
4JFD	ELAAIGILTV	10	HLA-A*02:01	2.46
4JFE	ELAGIGALTV	10	HLA-A*02:01	2.70
4JFF	ELAGIGILTV	10	HLA-A*02:01	2.43
4L3E	ELAGIGILTV	10	HLA-A*02:01	2.56
4MNQ	ILAKFLHWL	9	HLA-A*02:01	2.74
4OZG	APQPELPYPQPG	12	HLA-DQ (class II)	3.00
4P23	FEAQKAKANKAVD	13	H-2A (mouse, class II)	2.25
4P5T	FEAQKAKANKAVD	13	H-2A (mouse, class II)	3.26
5E9D	ELAGIGILTV	10	HLA-A*02:01	2.51
6AM5	SMLGIGIVPV	10	HLA-A*02:01	2.39
6AMU	MMWDRGLGMM	10	HLA-A*02:01	2.15

Table S1: **The fifty-structure panel.** PDB accession, native peptide, peptide length, MHC allele, and crystallographic resolution for every structure in the panel. The rule separates the twenty class I HLA-A*02:01-dominated structures above from the thirty added subsequently, which extend the panel to HLA-B, mouse H-2 and MHC class II.

Model	Params	Architecture	Training data	Predicted structures
ProteinMPNN	1.66M	Message-passing GNN	PDB (pdb_2021aug02), general backbones	–
ProteinMPNN (no MHC)	1.66M	Message-passing GNN (identical to above)	Same corpus, MHC/TCR-class structures excluded	–
ESM-IF1	141.7M	GVP-GNN + transformer	CATH-derived backbones	+12M structures AF2
LigandMPNN	2.62M	Message-passing GNN + ligand context	PDB, augmented with explicit ligand/heteroatom context	–

Table S2: **Models compared.** Parameter counts verified directly from each checkpoint’s `state_dict`. ESM-IF1 is the only tool whose training set is built substantially from AlphaFold2-predicted structures [15] rather than solved crystal structures (~12M predicted backbones, on top of its CATH-derived crystal set [29]).




Figure S1a: **Peptide length distribution** across all 50 structures in the panel.




Figure S1b: **Crystallographic resolution per structure.**




Figure S1c: **MHC residue vs. peptide contact density**, pooled across all 50 structures.




Figure S2: **Extreme redundancy at $T=0.1$.** 10,000 raw draws per (structure, model); the unique fraction is 0.50% on average and never exceeds a few percent.

Pooled design sequence logos, per model, by length -- pMHC+TCR
(4MJI/8-mer excluded -- singleton length, not poolable)




Figure S3: Design sequence logos, pooled by peptide length, one column per model.




Figure S4a: Per-structure design logos, ProteinMPNN. All 50 structures, ordered by native peptide then resolution.

Per-structure design enrichment logos (bits), ordered by peptide then resolution -- ProteinMPNN (no MHC)




Figure S4b: Per-structure design logos, ProteinMPNN (no MHC).

Per-structure design enrichment logos (bits), ordered by peptide then resolution -- ESM-IF1




Figure S4c: Per-structure design logos, ESM-IF1.




Figure S4d: Per-structure design logos, LigandMPNN.

Unique designs per model (mean $\pm$ SD across 50 PDBs; black dots = individual PDBs)
bracket = paired Wilcoxon, pMHC+TCR vs. pMHC only, n=50




Figure S5: Unique designs per model, mean $\pm$ SD across structures, pMHC+TCR vs. pMHC only.




Figure S7: **Resolution dependence in a controlled, same-peptide replicate set.** Four structures (2P5W, 2BNR, 2F53, 2F54), identical native peptide, near-identical TCR clone; every model plotted separately.




Figure S8a: ESMCBA vs. MHCflurry agreement, pMHC-only designs, by model.




Figure S8b: Per-model resolution/design-count relationship, full panel.




Figure S9a: **Each model's own confidence score for its designs shifts when the TCR is removed**, pooled across all forty HLA-A*02:01 structures. ProteinMPNN and ProteinMPNN (no MHC): average per-residue negative log-likelihood, lower is more confident. LigandMPNN: overall confidence, higher is more confident. All three shift toward lower self-assessed confidence without the TCR (Mann-Whitney $p < 10^{-200}$ in all three cases).




Figure S9b: **The own-score shift is not uniform across the panel.** Bonferroni-corrected significance (full vs. mhconly) per structure per model; dark cells are structures where the model's own-score distribution did not shift.




Figure S10: **Removing the TCR raises predicted affinity for every model, by very different amounts.** Strong-binder rate (MHCflurry IC₅₀ < 500 nM) per model in both structural arms; solid bars are pMHC+TCR, hatched bars pMHC only. Brackets give a Fisher's exact test of that model's own shift between arms, Bonferroni-corrected across the four models. Every model shifts up significantly, but the magnitude separates them: ProteinMPNN (no MHC) gains 25.7 percentage points and LigandMPNN 10.5, against 5.3 for ProteinMPNN and 4.0 for ESM-IF1. ESM-IF1 is the lowest-affinity model in both arms. The model ranking is preserved between arms, so the omnibus difference between models (reported above each arm) is not an artifact of which context is used.




Figure S11: **The sampled sequence space, colored by predicted binding affinity.** One UMAP (ESMCBA embeddings, fit jointly over every unique design), recolored six ways. All four models occupy the same overlapping landscape rather than four separate clusters, and MHCflurry and ESMCBA, built differently and trained differently, trace the identical gradient across it.




Figure S12: **Both MHC anchors are more hydrophobic than the peptide interior they sit alongside.** Kyte-Doolittle hydrophobicity of the native residue at P2 and P Ω for each of the 40 HLA-A*02:01 structures, plotted against that same peptide's mean over its own interior positions (P3 through the second-to-last, excluding P1). Grey lines connect the three values belonging to one peptide; anchor points are colored by side-chain class. Both anchors sit well above their own interior (paired Wilcoxon, $p < 0.0001$ for both, $n = 40$), so the two pockets do not differ in whether they prefer hydrophobic residues. They differ in how strictly the designs honor that preference, which is the distinction reported in Section 3.2. The points at P Ω colored as polar are chiefly the cysteine-terminated SLLMWITQC/NY-ESO-1 replicates, a residue whose discrete classification understates its hydrophobicity.

10-mers -- pMHC+TCR, $n=11$ structures $\times$ 4 models (violin = distribution shape, dots = every raw point; labels = % of points at $\geq 99\%$ and at $\leq 1\%$ recovery)




Figure S13: **Per-position recovery, decamers.** As Figure 1, for the 11 crystals with a 10-residue epitope. The P2 anchor again recovers above the interior, and the C-terminal anchor P Ω falls to 0.08, lower than in the nonamers.




Figure S14: **Per-position recovery, 13-mers.** As Figure 1, for the 5 crystals with a 13-residue epitope. Four of the five are MHC class II complexes, which have no P2 or P Ω anchor pocket, and the anchor structure seen at class I is correspondingly absent.




Figure S15: **Measured binding effect vs. recovery, pooled over every SKEMPI complex with designs.** Each point is one (complex, peptide position) with a measured single-mutation $\Delta\Delta G$; positive $\Delta\Delta G$ means the mutation weakens binding, marking a position the interface depends on. Left, recovery with the TCR present. Middle, the recovery gain from TCR context. Right, that gain against $\Delta\Delta G$ directly. 55 positions across 17 complexes, against the five positions in one crystal available before the panel was extended. $\Delta\Delta G$ marks positions the interface depends on, which includes MHC anchors as well as TCR contacts.

LigandMPNN score vs. published K_D , 9 DMF5 peptides with a measured value (4 excluded: K_D reported only as $>500 \mu\text{M}$ or not detected)




Figure S16: **The score carries no information about measured affinity.** LigandMPNN score against published K_D for the nine DMF5 peptides with an actual measured value, scored on 6AM5 (left) and 6AMU (right); points are colored by melting temperature and each structure's own crystallized peptide is circled. Four further peptides are excluded because their K_D is reported only as $> 500 \mu\text{M}$ or as not detected, which is an assay ceiling rather than a measurement. Neither backbone shows a relationship. Contrast this with Figure 4, where the same scores separate the two peptide families cleanly: these models place a peptide, they do not rank it.




Figure S17: **The anchor split is not an artefact of the panel's redundancy.** The panel holds 50 structures but only 33 distinct peptides and 26 distinct CDR3 β sequences, so a pooled statistic weights an epitope by how often it was crystallised. Left, P2, interior and PΩ recovery on progressively stricter sets: human MHC class I only, then one structure per TCR family (CDR3 β clustered so clonotypes differing by one or two substitutions collapse), then one per peptide, then restricted to HLA-A*02:01; each group is represented by its best-resolution member, chosen independently of any score. Right, the same result as the two gaps that define it. Both survive, and the PΩ deficit deepens from 0.23 to 0.35 as repeated epitopes are removed. The one exception is P2 on the TCR-only set, which at 17 structures is underpowered rather than contradictory.